Day 10: Part 1

Speakers: Prof Christian Richardt, University of Bath

Title: Towards Reconstructing + Editing the Visual World

Space life capture

4 Papers will be discussed in the sessions

Roadmap

#	Project	Venue
1	MatryODShka	ECCV 2020 (oral)
2	Deep Video Portraits	SIGGRAPH 2018
3	HoloGAN	ICCV 2019
4	BlockGAN	NeurIPS 2020

ECCV'20 ONLINE 23-28 AUGUST 2020

MatryODShka: Real-time 6DoF Video View Synthesis using Multi-Sphere Images

Benjamin Attal [1, 2] Selena Ling [1] Aaron Gokaslan [1] Christian Richardt [3] James Tompkin [1]





1 BROWN 2 Carnegie Mellon University 3 UNIVERSITY OF BATH

Multi-plane and multi-sphere images

- Multi-plane images
  - Perspective view synthesis
  - RGB + alpha layers
  - Inferred from plane sweep volumes

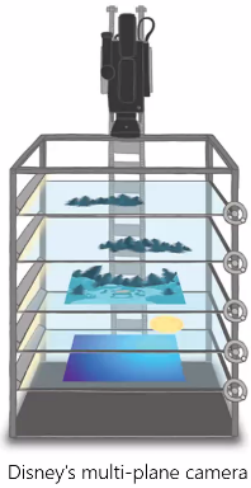

Disney's multi-plane camera

Multi-plane and multi-sphere images

- Multi-plane images
  - Perspective view synthesis
  - RGB + alpha layers
  - Inferred from plane sweep volumes

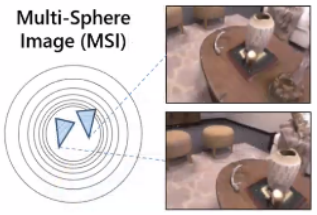

- Multi-sphere images
  - 360° view synthesis
  - RGB + alpha layers
  - Inferred from sphere sweep volumes

Matryoshka + ODS = MatryODShka

Multi-sphere image rendering

1. Intersect ray with each layer of MSI
2. Over-composite colors $\mathbf{c}_i$ and alphas α_i of intersection points:

$$\mathbf{c} = \sum_{i=1}^{N} \mathbf{c}_i \cdot \underbrace{\alpha_i \cdot \prod_{j=1}^{i-1}(1-\alpha_j)}_{\text{Net opacity of layer } i}$$

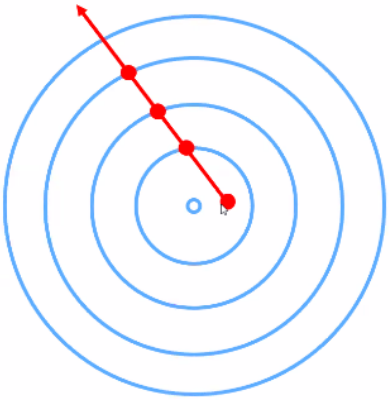

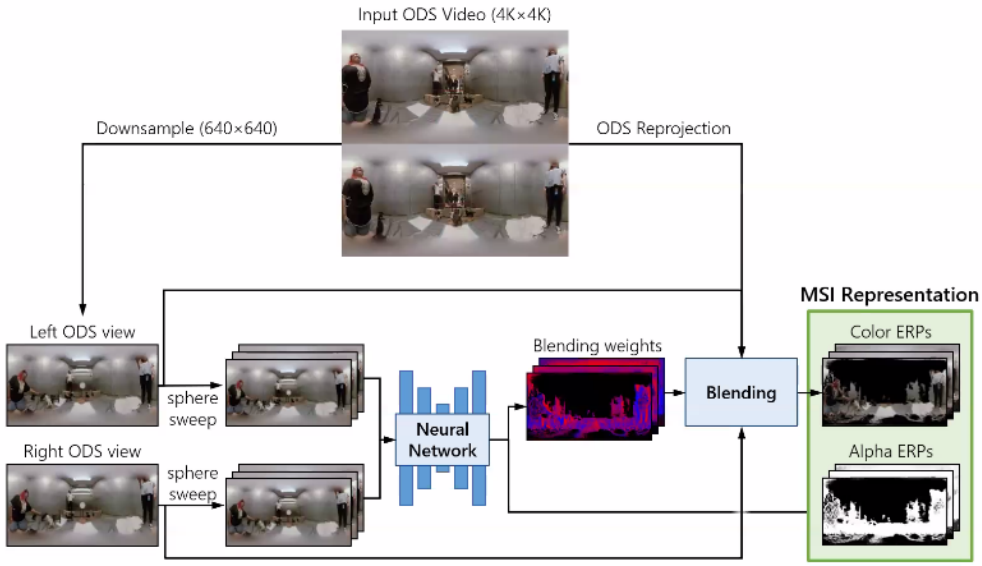

Spherically adapted training loss

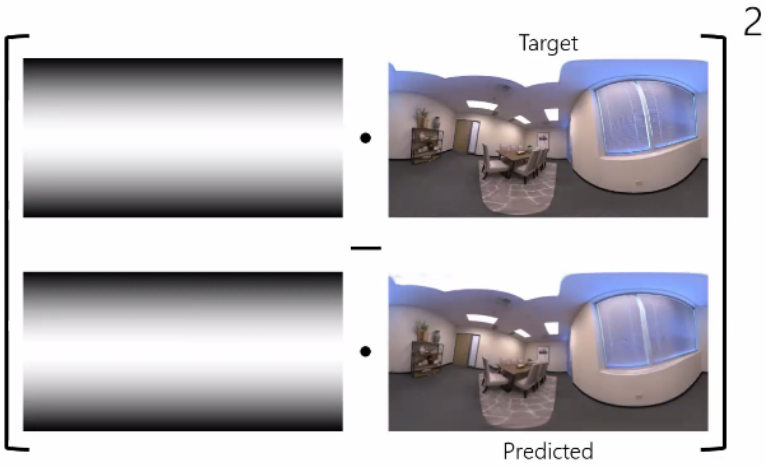

Training data

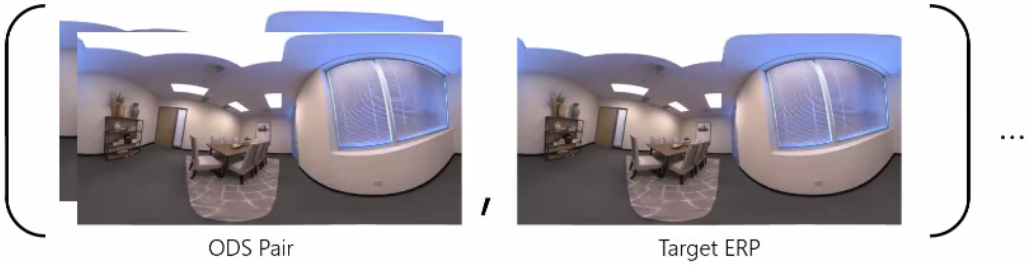

Straub et al. (2019). The Replica dataset: A digital replica of indoor spaces

Panoramic view synthesis

	Equirectangular View	Perspective View
Serrano et al. [IEEE VR 2019]		
MatryODShka [Attal et al. 2020]		

Concurrent work on synthesis

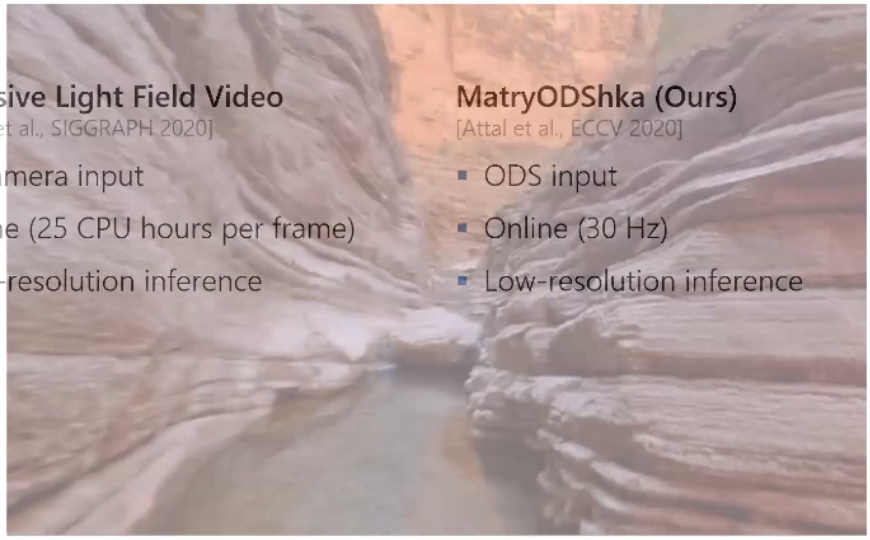

Immersive Light Field Video
[Broxton et al., SIGGRAPH 2020]

- 46 camera input
- Offline (25 CPU hours per frame)
- High-resolution inference

MatryODShka (Ours)
[Attal et al., ECCV 2020]

- ODS input
- Online (30 Hz)
- Low-resolution inference

Grand Canyon

Live view synthesis

Live view synthesis

Deep Video Portraits

Hyeongwoo Kim[1] Pablo Garrido[2] Ayush Tewari[1] Weipeng Xu[1]
Justus Thies[3] Matthias Nießner[3] Patrick Pérez[2]
Christian Richardt[4] Michael Zollhöfer[5] Christian Theobalt[1]

[1] MPI Informatics [2] Technicolor [3] TU Munich [4] University of Bath [5] Stanford University

GENERATIONS SIGGRAPH2018

Contribution

Original video Pose editing Expression editing

- Editing of head pose, rotation, face expression and eye gaze
- Combination of model-based face capture and CNN

Video: courtesy of UK government (Open Government Licence)

Head Pose along with expression Captured

Related Work

Model-based face capture and reenactment

Garrido et al., ToG 2016

Kemelmacher-Shlizerman et al., ECCV 2010
Shi et al., ToG 2014
Suwajanakorn et al., ICCV 2015
Thies et al., CVPR 2016
Averbuch-Elor et al., ToG 2017
Thies et al., SIGGRAPH 2018

CNN-based methods

Karras et al., ICLR 2018

Goodfellow et al., NIPS 2014
Isola et al., CVPR 2017
Chen and Koltun, ICCV 2017
Tewari et al., ICCV 2017
Olszewski et al., ICCV 2018
Wang et al., CVPR 2018

Overview

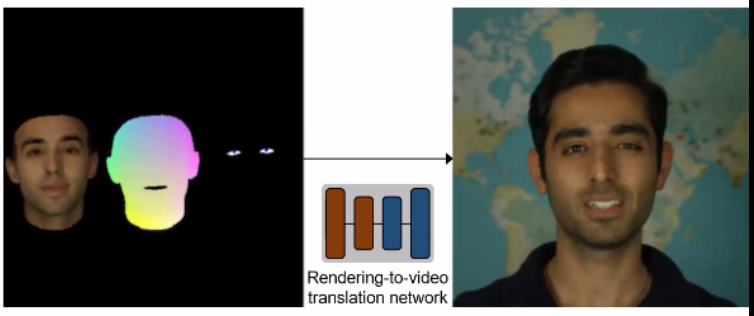

Synthetic rendering

Training video

Overview

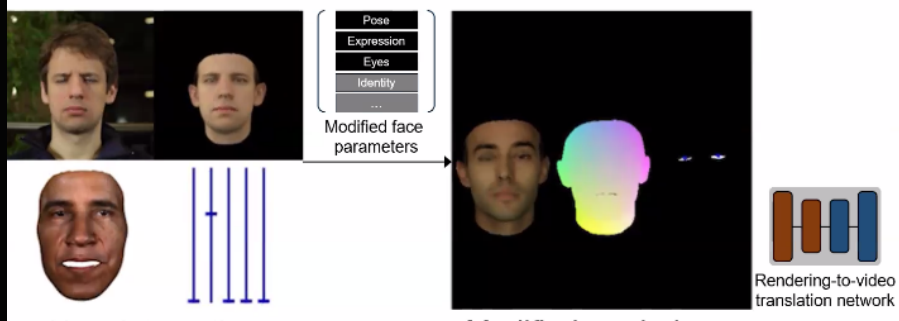

User interaction

Modified rendering

Overview

Rendering-to-Video Translation Network

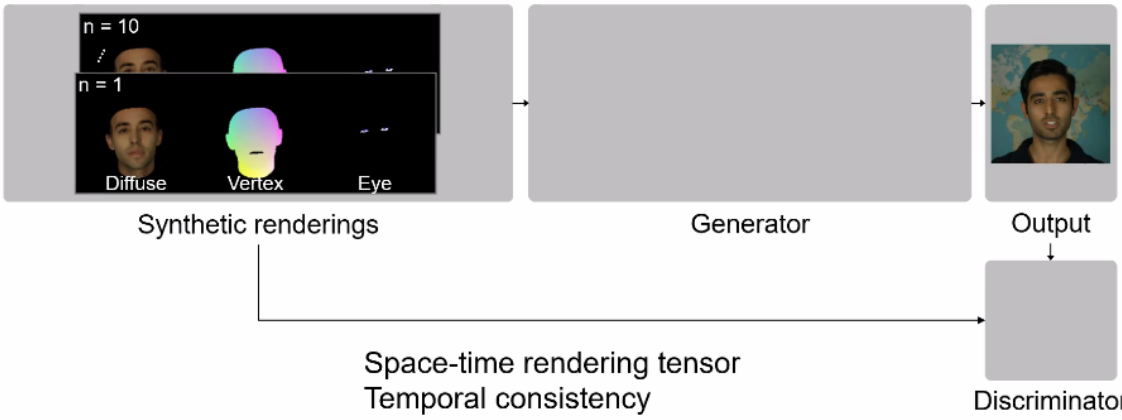

Rendering-to-Video Translation Network

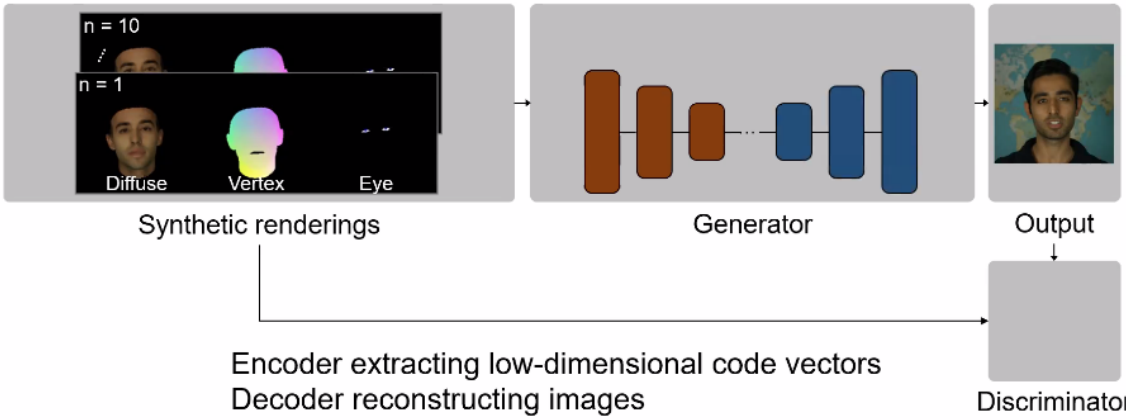

Rendering-to-Video Translation Network

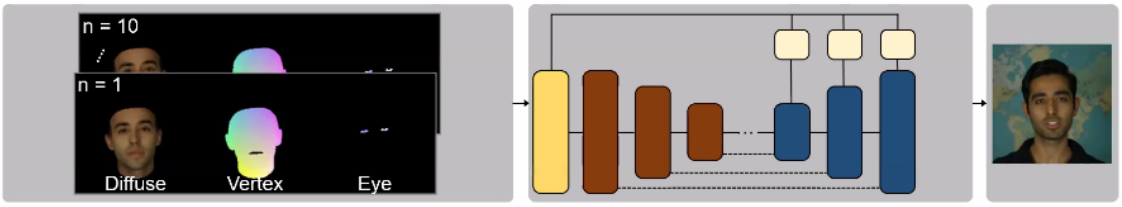

Synthetic renderings → Generator → Output → Discriminator

Skip-connections, multi-resolution and refinement
Fine-scale details

Chen and Koltun, ICCV 2017

Rendering-to-Video Translation Network

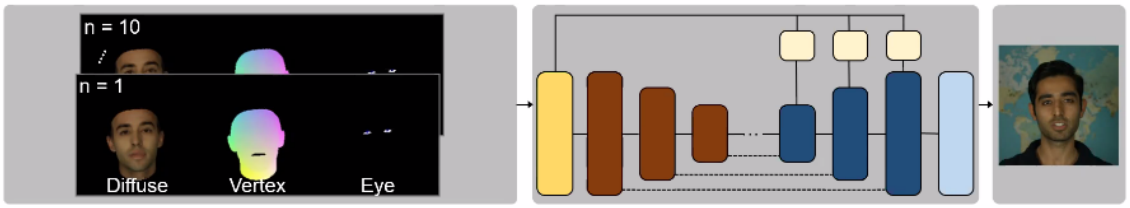

Synthetic renderings → Generator → Output → Discriminator

Training with adversarial and L_1 losses

Result: Facial Reenactment

Full reenactment of head pose, head rotation, face expression and eye gaze

Source Target

Result: Facial Reenactment

Full reenactment of head pose, head rotation, face expression and eye gaze

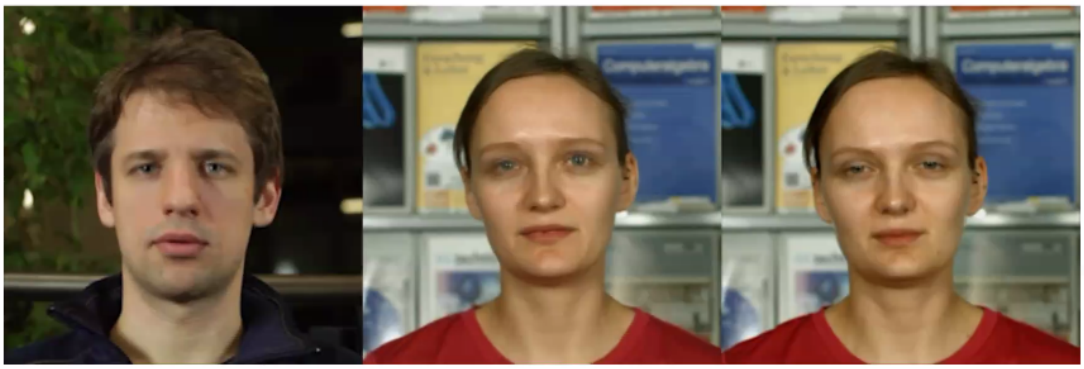

Source | Target | Face2Face (Thies et al., 2016)

Result: Facial Reenactment

Source | Target | Result

Video: courtesy of the White House (public domain)

Result: Visual Dubbing

Visual discomfort due to the discrepancy between video and audio tracks

Dubbing actor video | Original video

Result: Visual Dubbing

Modification of mouth motion to match audio tracks

Dubbing actor video | Dubbed video | Garrido et al., 2015

Result: Interactive Editing

Pose | Expression | Shape

Approximately 9 fps

Summary

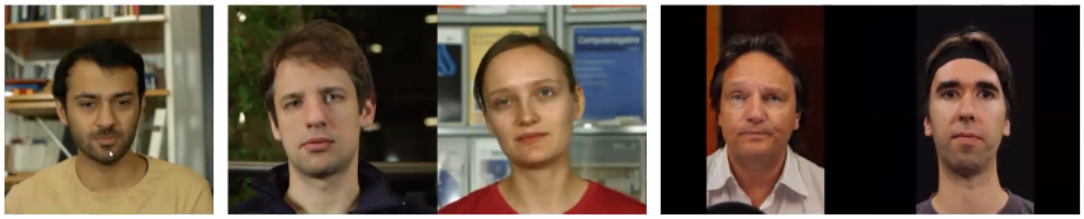

Future work:

- Pushing toward higher quality and resolution
- Video authentication and forensics

HoloGAN:
Unsupervised learning of 3D representations from natural images

ICCV 2019

Thu Nguyen-Phuoc Chuan Li Lucas Theis Christian Richardt Yong-Liang Yang

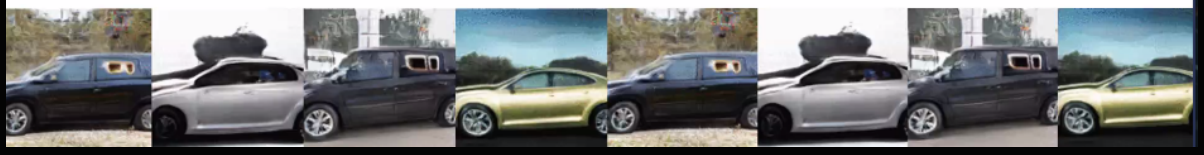

Generative adversarial networks

[Karras et al., StyleGAN, 2019]

*missed slide (to be updated)

Disentangled representation learning

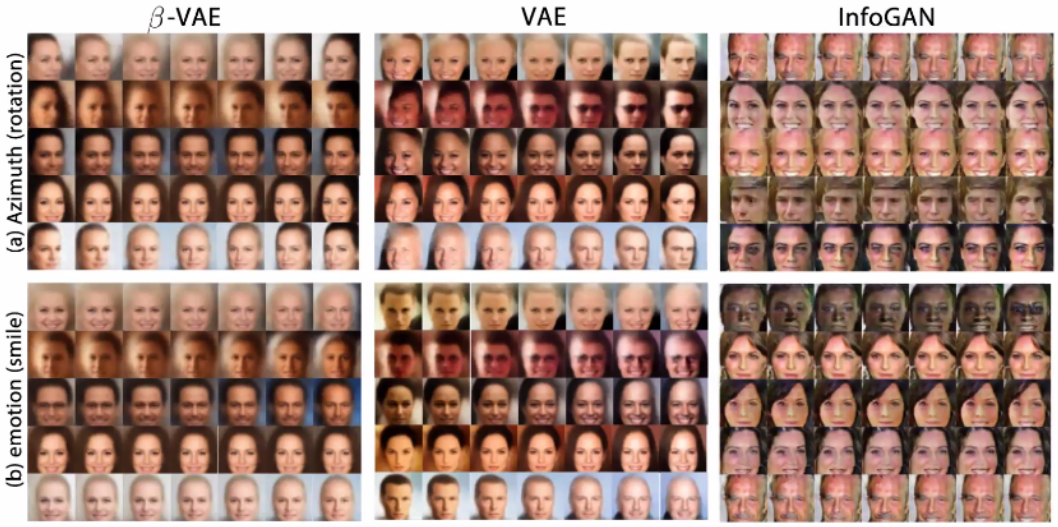

[Higgins et al., beta-VAE, 2014] [Chen et al., InfoGAN, 2015]

HoloGAN

Multi objects

Conditional GANs

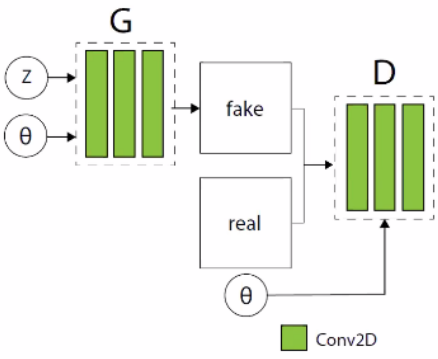

Computer Graphics Basics →

Back to the basics ...

Rendering comprises two main steps:

- Compute visibility
- Compute shading

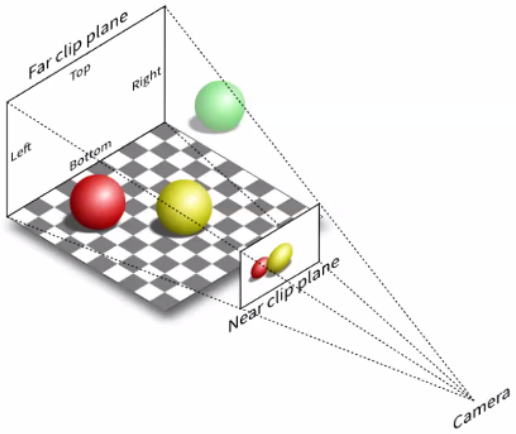

GAN architecture comparison

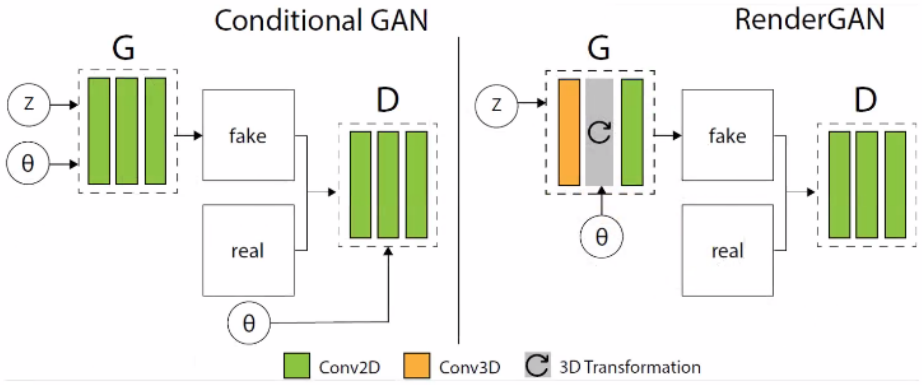

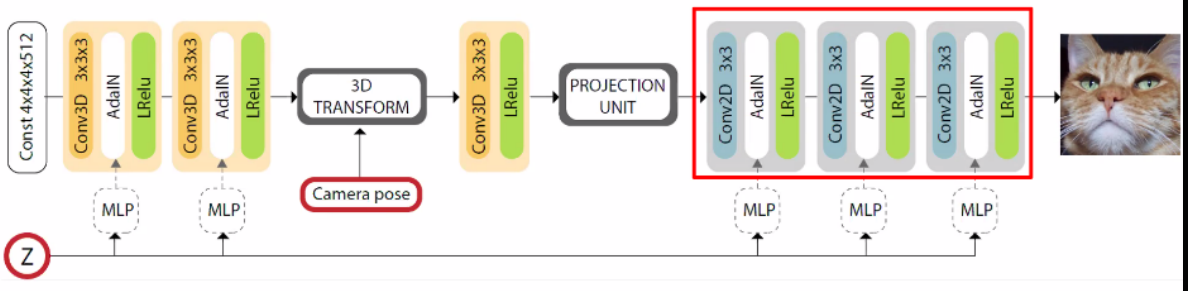

$$\min_G \max_D V(D,G) = \mathbb{E}_{x \sim p_{data}(x)}[\log D(x)] + \mathbb{E}_{z \sim p_z(z)}[\log(1 - D(G(T(z))))]$$

Datasets

Basel

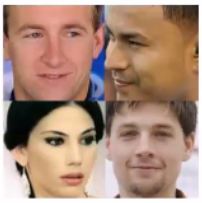

CelebA

Cats

Chairs

Cars

Cats

CelebA

Basel Face Model

Chairs

Cars

Style mixing

z1

z2

Separating shape + appearance

Separating shape + appearance

The next frontier — complex scenes

HoloGAN conclusion

- Adding inductive bias about the 3D world to a neural network
  - Better image quality
  - Better 3D understanding
  - Completely unsupervised

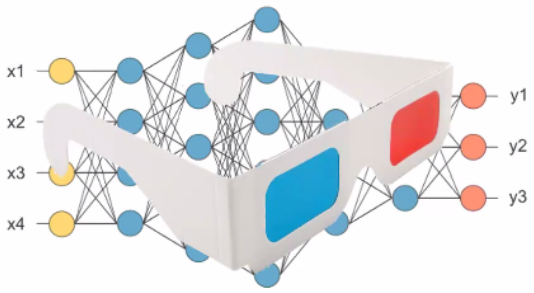

BlockGAN

Learning 3D object-aware scene representations from unlabelled images

NeurIPS 2020

Thu Nguyen-Phuoc Christian Richardt Long Mai Yong-Liang Yang Niloy Mitra

Current work

Scene Representation Networks
[Sitzmann et al. 2019]

HoloGAN
[Nguyen-Phuoc et al. 2019]

HoloGAN

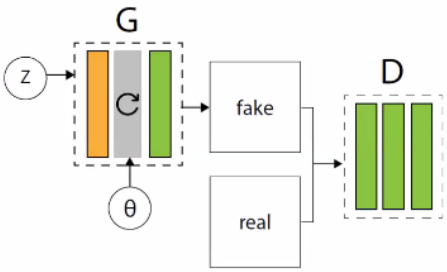

BlockGAN

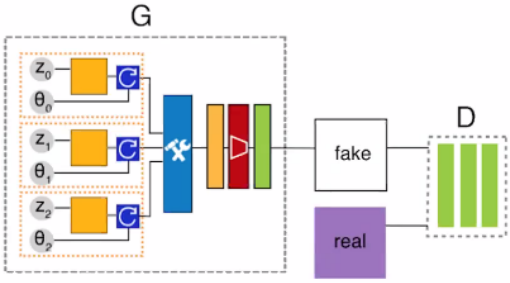

BlockGAN architecture

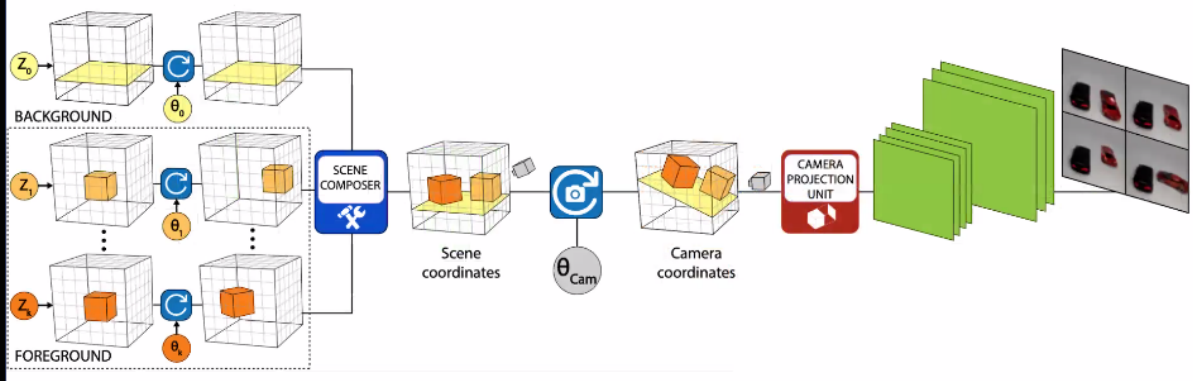

Datasets

Chair 1

Car 2

CLEVR 4

Cars (64×64)

Rotation

* all dataset rotated front & back

Translation

* Up & down translation

Changing background appearance

light & shadow

Changing object #1 appearance

Rotation

Translation

Change background

Change foreground

Comparison with 2D methods

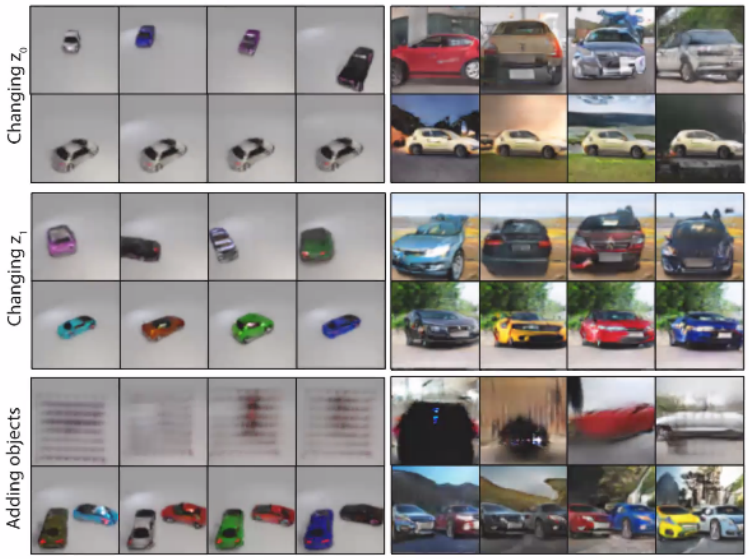

Generalisation

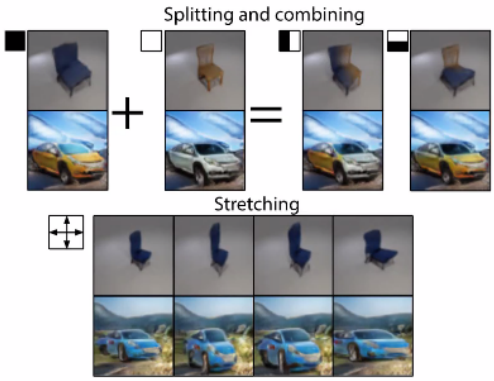

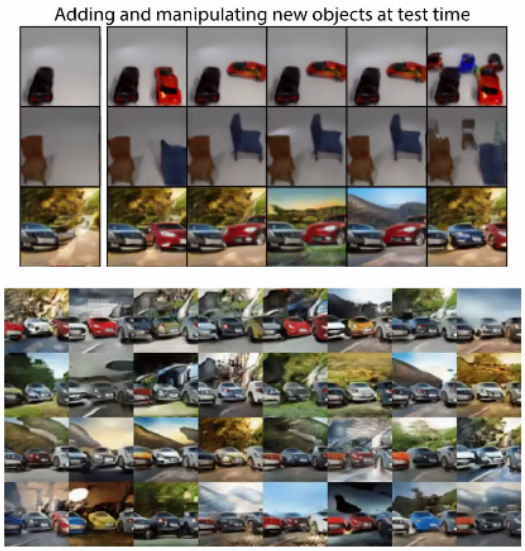

BlockGAN

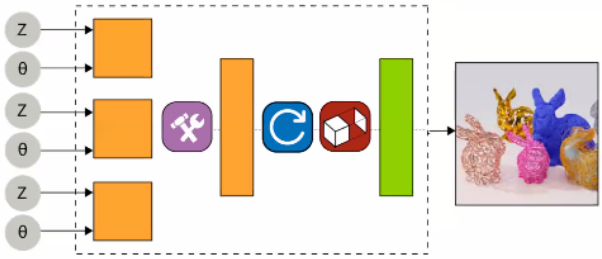

- BlockGAN offers control over pose and appearance of individual object in the scene
- Deep voxel representations offer intuitive object spatial manipulation and composition

Takeaways

- Traditional visual computing approaches build on explicit models
- Explicit models have limits:
  - sometimes difficult to model the real world, e.g. realistic faces
- Machine learning is reinvigorating visual computing:
  - can learn to model relationships from sufficient data (& expressive network)
  - no explicit model necessary, but often a black box …
  - deep learning in particular has shown amazing results
- Benefits to applying domain knowledge to structure networks:
  - helps learning + we know what's going on!

Horizon 2020
MSCA FIRE

Questions?

Christian Richardt

Towards reconstructing + editing the visual world

richardt.name
c_richardt